

Basanite Product Overview V1.0

1. Executive Summary

Basanite is an intelligent recruitment and interview system that transforms how organizations evaluate technical talent. By leveraging advanced AI capabilities, our platform generates customized interview frameworks that adapt to each candidate's background while maintaining consistent evaluation standards. This document provides a comprehensive overview of the problem we solve, our product design, user journeys, and our long-term aspirations.

2. Defining the Problem Space

2.1 The Measurement Gap

The current talent evaluation for high complexity roles in the tech industry are largely failing. Up until April 2026, these methods typically include Algorithmic Coding Tests, System Design Interviews, Structured Behavioral Interviews, Work Samples & Take-Home Projects, and Conversational Interviews. They are based on paradigms and frameworks invented a century ago, where human's ability to work was neither augmented nor weakened by their ability to leverage AI. On the surface, these methods appear objective and scalable. In practice, they select for the misleading proxy of interview preparedness, and are structurally blind to the qualities that actually drive performance in a world where AI handles an expanding share of routine technical execution.

With the AI revolution transforming talent needs at a fundamental level, this measurement gap presented by the current talent evaluation infrastructure is leading to both organizational and social consequences that compound as AI capabilities advance.

For corporations, hiring processes optimized for the wrong signals produce systematic mismatches between role requirements and the people placed into them. These are invisible at the point of hire, but manifest downstream in slower execution, higher management overhead, and accelerating attrition. Moreover, when evaluation frameworks cannot reliably identify who will drive product velocity, navigate ambiguity, or lead effective human-AI collaboration, organizations lose their ability to build the capability portfolios that competitive advantage depends on.

For society, as AI increasingly handles execution-layer tasks, the skills that generate human economic value are shifting toward judgment, synthesis, ethical reasoning, and collaborative intelligence. Yet talent pipelines continue to screen for the skills AI is replacing fastest, while filtering out the human qualities that remain irreplaceable. Over time, this produces a workforce allocation problem of significant scale, where individuals with genuinely differentiated capabilities are sorted into the wrong roles or excluded from high-value work entirely, while organizations fill critical positions with candidates who are optimized for a mode of work that is rapidly becoming obsolete.

We are therefore in urgent need of critical reassessment of talent evaluation theories, and an innovative reinvention of talent evaluation paradigms and tools.

2.2 Human Intelligence That Matters

Research spanning cognitive science, philosophy of knowledge, behavioral decision theory, and organizational psychology converges on a different account of what drives genuine value in technical organizations. The qualities that distinguish high performers in complex, adaptive, AI-era work environments are not the ones that conventional hiring instruments capture. They include:

- **Judgment under ambiguity:** the capacity to act decisively on incomplete information, without either paralysis or false confidence
- **The ability to surface and transmit tacit knowledge:** what Michael Polanyi described as the irreducible gap between what experts know and what they can say: the practical intelligence that lives in experience rather than in text
- **Sound intuition in the absence of sufficient data:** what Gary Klein's decades of naturalistic decision-making research revealed as the recognition-primed judgment that distinguishes genuine experts from those who merely know the vocabulary of expertise
- **Psychological safety and collective learning:** the social and epistemic conditions under which teams can correct their own errors before they compound
- **Creative problem reframing:** the ability to recognize when the team is solving the wrong problem; a capacity that is actively suppressed by evaluation systems that reward convergence on expected answers
- **Ethical reasoning:** not the ability to articulate ethical frameworks, but the capacity to feel the weight of real tradeoffs and navigate them with integrity
- **The capacity to be genuinely transformed by experience:** the metacognitive self-awareness that separates those who accumulate experience from those who learn from it

These qualities share a structural feature that explains why conventional hiring instruments cannot detect them: they cannot be retrieved from a knowledge base. They are not facts to be recalled or algorithms to be executed. They are capabilities forged through real

experience, expressed through real judgment, and legible only to evaluators who know what to look for and how to ask.

3.Design Philosophy

Assess Genuine Capability, Not Performed Competence

The central conviction behind every design decision at Basanite is simple: most hiring tools are optimised for the wrong signal. They measure how well someone can approximate the idea of a good candidate, rather than whether they actually are one.

3.1 Five Principles

3.1.1 Depth Over Breadth

Basanite is designed around progressive excavation: each layer of the assessment exists to move one level deeper into signal quality. A candidate who answers fluently at the surface should encounter ground that shifts beneath them at the next layer. The edges of real ability are blurry. Performed ability has no edges.

3.1.2 Structure as a Fairness Mechanism

Unstructured evaluation produces advantage for those who are most culturally legible to the evaluator. Standardized evaluation encourages cheating and rewards those who perform. In the absence of valid capability signals, assessors default to cultural familiarity. Candidates who mirror the background, communication style, and social register of the interviewer are advantaged not because they are more capable, but because they are easier to read.

Basanite's structure aim to be equalising. By anchoring every evaluation to consistent frameworks, principled questioning logic, and explicit scoring criteria grounded in observable evidence, the system creates conditions where a self-taught engineer without institutional pedigree can be seen as clearly as one from a target university, as both are being asked the same questions in the same spirit, with the same depth of follow-up.

3.1.3 The Encounter as Part of the Assessment

The experience a candidate has inside a Basanite evaluation should feel like a rigorous but genuinely curious conversation instead of an interrogation, a performance, or a test of how well they have prepared for exams. This matters for two reasons:

1. The collaborative, uncertainty-tolerant environments that the best technical organisations are building require people who can think under genuine pressure, not people who can perform competence in artificial conditions. An evaluation system that rewards the latter is selecting against the former.
2. Candidates who feel genuinely seen regardless of outcome become the most powerful distribution channel available. The reputation of an evaluation process travels through the talent communities it touches. A process that candidates respect is one that attracts the candidates worth reaching.

3.1.4 Honest About What AI Can and Cannot Do

Basanite will never claim more than the evidence supports. The system is strongest where AI has a genuine structural advantage over human evaluation: in consistency, in scale, in the ability to follow a thread across a long conversation without fatigue or drift, and in the capacity to maintain identical standards regardless of the order in which candidates are seen. These are real advantages, and they compound. The tenth candidate in a Basanite evaluation receives the same quality of assessment as the first. This is not true of human-led processes, and the difference matters at scale.

But there are things that require irreducible human judgment, such as nuances of seniority assessment, domain-specific technical verification, and the final hiring decision itself. Basanite is designed to be explicit about these limits: flagging where human expertise is required, producing quotable evidence rather than opaque scores, and positioning itself as the infrastructure that makes human judgment better rather than the mechanism that replaces it. Intellectual honesty about capability boundaries is a design requirement, not a marketing disclaimer.

3.1.5 Evaluation as a Two-Way Mirror

The best hiring processes leave candidates with a clearer understanding of themselves instead of just a verdict. A process that produces genuine signal about a candidate's capabilities should, by definition, produce signal that is legible to the candidate as well as the organisation. If a candidate cannot understand what was being assessed and why, the assessment itself is suspect.

This principle has a practical implication for system design: every evaluation strategy deployed by Basanite must be one that can be honestly explained to the candidate it is applied to. Strategies that depend on deception to function that work only because the candidate

does not know what is being measured, are excluded not merely on ethical grounds, but on epistemic ones. An evaluation that cannot withstand transparency is not producing the signal it claims to produce.

4. Prototype Architecture

5. Defining Target Customers

5.1 Primary Target Customers

Tier 1: High-Volume Technical Recruiters: Enterprise and Scale-Up

Any organisation running more than 30 to 40 technical hires per year has a structural headcount cost in screening and first-round assessment that Basanite directly displaces. This includes:

- High-growth technology companies at Series B and beyond with dedicated talent teams
- Technology divisions of large financial institutions, retailers, and professional services firms that hire engineers at volume but are not primarily technology companies
- Management consultancies and technology consultancies running graduate and experienced hire pipelines

For this customer, the primary commercial argument is not quality but headcount. One Basanite deployment can replace the screening and first-round assessment work of one to three full-time recruiters. At £60,000 to £90,000 per recruiter fully loaded, the payback period on a Basanite contract is measured in months, not years. Quality improvement is the reason the legal and ethical case for the product holds; cost displacement is the reason it gets purchased.

Tier 2: Staffing Agencies and Recruitment Process Outsourcing Firms

Staffing agencies and RPO firms face acute margin pressure as the market commoditises. Their differentiation, increasingly, must come from the quality and defensibility of the assessment they provide instead of only from their candidate databases and client relationships. A

Basanite-powered agency can offer enterprise clients a materially stronger assessment product at lower per-candidate cost than a human-led screening process, while improving their own gross margin by reducing the recruiter hours per placement.

The incentive misalignment exists for agencies whose business model depends on volume throughput of weak candidates. It does not exist for agencies attempting to move upmarket toward retained search and quality-differentiated assessment services. The latter is a growing segment, and Basanite is a natural infrastructure partner for that strategic shift.

Tier 3: SMEs Without Dedicated HR Functions

Small and medium-sized technology companies, typically between 20 and 150 employees, often have no dedicated HR or recruiting function at all. Hiring is done directly by founders, CTOs, or engineering leads who do not have the bandwidth for structured screening and lack the institutional knowledge to design rigorous first-round assessments. These companies consistently over-rely on personal networks and referrals, which produces the homogeneous teams that the research literature predicts: brittle, culturally narrow, and poorly calibrated for the uncertainty of early-stage growth.

For this customer, Basanite is not replacing existing headcount, it is providing a capability that does not exist. The commercial framing is closer to infrastructure than to labour substitution: Basanite gives a 30-person company the evaluation sophistication of a company ten times its size, at a fraction of the cost of hiring a Head of Talent to build that sophistication internally. This segment is price-sensitive but large, and the switching cost once embedded in a hiring workflow is substantial.

Tier 4: Universities, Graduate Employers, and Professional Training Programmes

As previously identified, this segment provides both revenue and strategic data value. The revised commercial argument here is more pointed: universities and large graduate employers are running assessment centres that cost £500 to £2,000 per candidate in assessor time, venue costs, and coordination overhead. For a graduate scheme processing 2,000 applications down to 50 offers, the assessment infrastructure cost alone can exceed £500,000 per cycle. Basanite can execute the first two to three stages of that funnel at a fraction of the cost, with demonstrably superior consistency.

Tier 5: Geographically Distributed and Remote-First Organisations

A structural shift in how technical work is organised has created a specific customer need that was not present five years ago. Organisations hiring across multiple time zones and jurisdictions face a coordination cost in scheduling and running human-led assessments that compounds with geographic distribution. Asynchronous AI-led evaluation removes this friction entirely. For a company

hiring engineers across three continents from a central talent function, the operational value of an always-available, timezone-agnostic assessment system is significant and immediate.

5.2 Who Basanite Is Not For At This Stage

Non-technical roles, at least initially. Basanite's assessment architecture is purpose-built for technical capability evaluation. The question design logic, the tacit knowledge detection framework, the cheating identification signals, and the transfer testing methodology all assume a candidate with a technical background being evaluated for a technical role. Extending to sales, operations, legal, or marketing roles is not impossible — but it would require a substantive rebuild of the core assessment logic, not a surface-level adaptation. This is a deliberate scope constraint, not a permanent exclusion.

Roles where the evaluation is primarily relational or political. Senior executive hiring — C-suite, VP-level general management, board appointments — is governed by dynamics that no structured assessment system currently handles well. These decisions involve stakeholder alignment, organisational politics, and relationship capital that are genuinely resistant to systematic evaluation. Human judgment, executive search relationships, and reference networks remain the dominant signal sources in this tier. Basanite may eventually have a role here, but it is not the right first market.

Organisations that have already decided they want a human face on every candidate interaction. Some companies, particularly those that market heavily on employer brand and candidate experience, have made a deliberate choice that every touchpoint in their hiring process will be human-led. This is a values position, not a capability gap, and it is not one that Basanite should attempt to argue against. The organisations that hold this position most strongly tend to be in consumer-facing industries where brand perception is tightly managed. They are not the right early adopters.

6. User Journeys

6.1 The Hirer's Journey

Stage 1: Account Setup and Job Configuration

The hirer logs into their Basanite account and navigates to the job creation flow. They paste or compose a job description in the same format they would use on LinkedIn, Greenhouse, or any standard hiring platform: role title, responsibilities, requirements, and any contextual information about the team or company. No additional formatting or translation is required. Basanite accepts the job description as-is.

Basanite analyses the job description and surfaces a recommended set of evaluation dimensions drawn from its assessment library. The recommendation is generated by matching the role's requirements, seniority signals, and technical domain against Basanite's dimension taxonomy. For a senior backend engineer role, for example, the system might recommend dimensions including technical judgment depth, judgment under ambiguity, and the ability to surface tacit knowledge; for a junior data analyst role, it might weight intuition under data scarcity and creative problem reframing differently. The hirer sees the recommended combination with a brief rationale for each dimension.

The hirer can accept the recommendation as a complete package, remove dimensions they consider less relevant to this specific role, or browse the full dimension library and add dimensions not included in the default recommendation. The interface presents dimensions with plain-language descriptions and example signal indicators, so the hirer does not need to understand the underlying assessment methodology to make informed choices.

Stage 2: Eligibility Constraints and Budget

Once the evaluation dimensions are confirmed, Basanite prompts the hirer to specify any hard eligibility constraints that should gate candidates before assessment begins. These include:

- Right to work in a specific jurisdiction
- Graduate or current student status
- Minimum or maximum years of professional experience
- Specific technical prerequisites and how negotiable they are for the role
- Location or timezone requirements

These constraints are stored as pre-screening filters. Candidates who do not meet them are redirected without consuming assessment capacity, and receive a respectful notification that they do not meet the eligibility criteria for this specific role at this time.

Basanite then asks the hirer to specify their assessment budget, either as a total spend for the hiring round or as a per-candidate cap. This budget determines the depth and scope of assessment each candidate receives: the number of evaluation dimensions actively probed, the

length of the interview, and whether advanced features such as transfer testing or knowledge reproduction assessment are activated. The system presents the hirer with a clear mapping between budget tiers and assessment scope before they confirm.

Stage 3: Role-Specific Prompt Generation

With job description, evaluation dimensions, eligibility constraints, and budget confirmed, Basanite generates a role-specific interview agent prompt. This prompt is derived from Basanite's master assessment framework and calibrated to the specific parameters the hirer has set: the target dimensions, the technical depth tier appropriate to the role's seniority, and the candidate experience path determined by whether candidates are likely to have directly relevant prior experience.

The generated prompt is stored against the role and activates automatically when candidates begin their assessment. The hirer does not see or interact with the prompt directly, it is operational infrastructure, not a configurable document. The hirer can preview the types of questions and assessment areas the interview will cover in plain language, but the specific prompt logic remains internal to protect assessment integrity.

The role is now live and ready to receive candidates.

6.2 The Candidate's Journey

Stage 1: Application and Onboarding

The candidate arrives at the Basanite assessment link, either directly from a job posting or via an invitation from the hirer. They create a Basanite account or log into an existing one. If they have completed Basanite assessments for other roles previously, their profile history is retained, but each assessment is conducted fresh against the specific role's parameters.

The candidate uploads their CV or résumé. Basanite processes this document and extracts the candidate's experience history, technical background, project work, and any signals relevant to the role being assessed. This information is used to personalise the assessment encounter. The system identifies which of the candidate's experiences are most relevant to the target role and uses these as anchor points for the interview.

Before the assessment begins, Basanite presents the candidate with a brief plain-language explanation of what the assessment involves, approximately how long it will take, and what kinds of questions they can expect. The tone is transparent and respectful: candidates are

told they will be asked about their real experiences and their genuine ways of thinking, and that there are no trick questions or expected answers they should be trying to perform toward.

Stage 2: Assessment

The candidate enters the interview. Basanite begins with an open invitation for the candidate to walk through a recent piece of work in their own words: a project, a decision, a problem they solved. This narrative foundation serves multiple functions: it establishes a shared factual basis for follow-up questions, it gives the candidate a comfortable entry point, and it gives the system an initial profile of the candidate's experience depth and communication style against which to calibrate the remainder of the assessment.

From this foundation, Basanite moves through the active evaluation dimensions confirmed for the role. Questions grow organically from the candidate's own narrative rather than from a fixed script. Where the candidate has demonstrated relevant prior experience, the system deploys real-experience questions that probe the depth and authenticity of that experience. Where relevant experience is absent or thin, the system activates open sandbox questions designed to surface the candidate's reasoning and judgment in hypothetical but realistic scenarios.

Throughout the interview, Basanite tracks narrative consistency, monitors for signals that distinguish genuine capability from performance, and calibrates follow-up depth based on the signal confidence it is accumulating in each evaluation dimension. The interview continues until signal confidence thresholds are met across the required dimensions, or until the budget-determined time limit is reached.

The candidate experience throughout is designed to feel like a substantive, curious conversation. Basanite does not rush, does not deploy trick questions, and does not penalise candidates for acknowledging uncertainty. Candidates who say "I don't know" in a way that demonstrates genuine metacognitive awareness are treated differently from candidates who produce confident but hollow answers — because the former is a positive signal, not a failure.

Stage 3: Dual Feedback Reports

When the assessment concludes, Basanite generates two separate feedback documents.

The hirer report is a structured assessment output containing dimension-by-dimension scores, the specific candidate statements that support each score, flagged areas of particular strength or concern, a cheating risk assessment where applicable, a technical capability map identifying areas of depth and potential blind spots, and a one-paragraph synthesis of what this candidate is likely to be like in a real

working environment. This report is designed to function as a briefing document for the hirer's final human-led interview: it tells the interviewer where to probe further, what has already been established, and what remains genuinely uncertain. It does not make a hire or no-hire recommendation, as the hiring decision remains with the human.

The candidate report is a brief, neutral summary of their performance delivered in plain language. It acknowledges what they demonstrated well, identifies areas where they could strengthen their presentation or preparation, and offers constructive suggestions for development. The candidate report is deliberately designed to be informative without being reverse-engineerable: it does not reveal the specific evaluation dimensions used, the scoring criteria, or the question logic underlying the assessment. A candidate cannot use their feedback to game a subsequent Basanite assessment for the same or a different role. The tone is respectful and genuinely useful, a candidate who did not progress should come away with something of value from the process.

Stage 4: Ranking and Queue Management

Once the candidate's report is generated, Basanite's ranking algorithm incorporates the assessment output into a composite score for this candidate against this role. The composite draws on dimension scores, background signals extracted from the CV, stated salary expectations, and any hard eligibility data captured during onboarding.

The candidate is assigned a position in the hirer's live ranking queue for this role. The ranking is dynamic: as new candidates complete assessments, the queue updates automatically. The hirer sees a continuously ordered list of assessed candidates, each with their composite score, a summary of their strongest and weakest dimensions, and a direct link to their full hirer report.

The hirer can sort and filter the queue by individual dimensions, salary band, experience level, or composite score. When they are ready to move candidates forward to a final human-led interview, they select from the queue with the context of Basanite's full assessment in hand, knowing not just where each candidate ranked, but why, and what the human interview should focus on to complete the evaluation.

7. Technical Innovation and Differentiation

7.1 Basanite Interview Technique Inventory

7.1.1. Structural Techniques

Macro Three-Part Structure

Every interview follows a fixed high-level arc: narrative foundation → dimension assessment → stress test. Candidates must establish their own story before probing begins.

Signal Saturation Pacing

Transitions between dimensions are governed by internal confidence thresholds, not question counts. The interviewer moves on when it has enough, not when a quota is met.

Experience Path Bifurcation

The interview forks at the outset based on whether the candidate has relevant prior experience. Path A deploys experience-anchored questions; Path B deploys capability-inference questions with a weighted composite scoring model.

Technical Depth Calibration

Before the interview begins, the agent calibrates to one of two depth registers: applied or research/architecture, based on the job description, which governs the intensity of technical probing throughout.

7.1.2. Questioning Techniques

Narrative Anchoring

The opening phase invites the candidate to describe a recent piece of work in their own words, establishing a shared factual baseline from which all subsequent questions grow organically.

Experience-Grounded Questioning

Questions are tied to the candidate's own stated history rather than hypothetical scenarios. The container is always the candidate's real past; the question only works if they actually did the thing.

Open Sandbox Questioning

Where real experience is absent or insufficient for a particular evaluation metrics, the candidate is placed in a realistic hypothetical scenario. Five specialist sandbox design principles govern these questions to prevent them from being answerable by AI lookup alone.

Decision Tracing

The candidate is asked not only what they chose but why they did not choose the alternatives. The quality of their dismissal of Option B

reveals more than their justification of Option A.

Boundary Condition Probing

Rather than asking whether a solution works, the interviewer asks under what conditions it fails. Valid technical knowledge has edges; performed knowledge does not.

Counterfactual Pressure

In the stress test phase, the candidate is asked what would have happened if they had made the opposite decision. This requires sustained reasoning in a direction for which they have no prepared answer.

Forced Commitment

When a candidate gives a framework-level answer, the interviewer demands a specific, defensible position. This disrupts the AI tendency toward balanced non-answers.

Compulsory Trade-Off Framing

Sandbox questions are structured to make "consider all dimensions" a wrong answer strategy. Resource constraints or mutually exclusive choices are built into the scenario to force prioritisation.

7.1.3 Follow-Up and Depth Techniques

Progressive Excavation

Each layer of follow-up is designed to go one level deeper into the same territory rather than expand to new ground. The goal is vertical depth, not horizontal coverage.

Vagueness Targeting

When a candidate gives an answer that sounds complete, the interviewer identifies the single most ambiguous detail within it and asks the candidate to elaborate specifically on that point.

Fourth Wall Break

Mid-scenario, the interviewer pulls the candidate out of the hypothetical and asks what was genuinely the first thing that came to mind, not the answer they organised, but the raw initial reaction. AI has no first reaction; humans do.

Honest Failure Elicitation

The candidate is asked to name a decision in their work that they did not think through fully but that happened to turn out fine anyway. Clean answers to this question are a high-risk signal. Slightly uncomfortable, specific ones are not.

Predict-Your-Own-Error

After a candidate gives a judgment, they are asked where that judgment is most likely to be wrong, not theoretically, but based on their own known patterns of error. Generic risk lists are a signal; personal blind spots are the target.

Regret Probing

For any experience the candidate describes as a success, the interviewer asks what they would have done differently. Authentic experience always contains regret; performed experience does not.

7.1.4 Consistency and Authenticity Verification Techniques

Narrative Consistency Tracking

All specific details mentioned by the candidate: team sizes, timelines, technical parameters, project names, are internally logged and referenced in later questions to test whether the account holds together across the full conversation.

Parameter Verification

For technical claims, the interviewer asks for concrete numbers: dataset sizes, latency figures, team headcount, timeline specifics. Real memory is approximate and slightly uncertain; fabricated or AI-assisted memory tends to be either implausibly precise or implausibly round.

Cognitive Priority Testing

A key detail is buried in a scenario description, one that an experienced practitioner would immediately identify as the most important signal, but that an inexperienced candidate or AI would treat as background noise. The candidate is assessed not on what they answer but on what they notice.

Information Gap Injection

A critical variable in a scenario is deliberately left ambiguous, one that only real experience can resolve through assumption. A candidate with genuine experience names a specific calibration factor; one without produces a generic framework.

7.1.5 Anti-Cheating Techniques

Structural Opacity

Evaluation dimensions are never disclosed to the candidate. Question purposes are concealed. The candidate cannot reverse-engineer what is being measured and optimise their performance toward it.

AI Output Signal Detection

The interviewer maintains a checklist of linguistic and structural patterns associated with AI-generated responses: comprehensive coverage without prioritisation, clean narrative arcs without friction, self-criticism that reads as a growth story, absence of any moment of genuine uncertainty.

Cognitive Load Escalation

The deepest follow-up is reserved for the areas where the candidate is most confident, not most vulnerable. Defences are lowest where people feel they are winning; that is where authentic signal is easiest to surface.

Latency Awareness

Anomalous response timing: answers that arrive too quickly for genuine reflection or too slowly for normal recall, is flagged as a risk indicator.

Tacit Knowledge Consistency Test

A candidate who claims a skill they cannot teach, but then describes its internal mechanism with unusual precision, has produced a self-contradicting account. The inability to articulate is part of the claim; fluent articulation undermines it.

7.1.6 Scoring and Output Techniques

Quotation-Anchored Scoring

No dimension may receive a score above three without a specific verbatim statement from the candidate cited as evidence. Scores are grounded in what was actually said, not in overall impression.

Cheating Risk as Independent Dimension

Cheating risk is scored and reported separately from capability. A candidate who may have used AI assistance but produced strong responses is treated differently from one who produced weak responses. The two assessments do not collapse into each other.

Dual Report Generation

Two separate outputs are produced from the same assessment: a hirer report designed to brief the human final interviewer, and a candidate report designed to be useful without being reverse-engineerable. The same underlying data serves two different communicative purposes.

Technical Capability Mapping

Rather than a single technical score, the output produces a structured map: areas of demonstrated depth, areas of surface fluency without underlying understanding, specific nodes requiring human technical expert verification, and surface-fluent but substantively suspect domains.

8. Academic Partnerships

9. Corporation Partnerships

10. Validation Plans