Basanite Product Overview V2.1

> **Revision summary relative to V1.0.**
> (i) Sharpened the problem statement around the failure of the *technical* interview, specifically coding-assessment cheating, and the unmeasured AI-collaboration ability of engineers.
> (ii) Added an explicit scope statement: Basanite addresses the technical layer of the interview only, not the psychometric or culture-fit layers.
> (iii) Added a Tech-Industry Map section enumerating the verticals, roles, and seniority bands Basanite calibrates for.
> (iv) Rewrote the metacognitive-dimensions section with formal definitions, expanded to eight dimensions, and added references to the underlying literature.
> (v) Replaced the single-round assessment methodology with a two-round architecture, with a new Round 2 (AI Collaboration Workbench) modelled on openround.ai's published approach. All other sections of V1.0 are preserved verbatim.

Table of Contents

1. Executive Summary

Basanite is an intelligent recruitment and interview system that transforms how organizations evaluate technical talent. By leveraging advanced AI capabilities, our platform generates customized interview frameworks that adapt to each candidate's background while maintaining consistent evaluation standards. This document provides a comprehensive overview of the problem we solve, our product design, user journeys, and our long-term aspirations.

---

2. Defining the Problem Space

2.1 The Failure of the Technical Interview

The talent-evaluation infrastructure for high-complexity roles in the tech industry is no longer fit for purpose. The breakdown is most acute in the *technical* portion of the interview, which is the part that is supposed to verify whether a candidate can actually do the work. As of April 2026, the dominant instruments including algorithmic coding tests, system-design interviews, technical phone screens, take-home projects, and structured technical conversations, were designed for a world in which a human's unaided cognition was the primary determinant of engineering output. That world no longer exists.

Two specific dysfunctions have crystallized over the past 24 months, and together they hollow out the technical interview as a measurement instrument.

Coding assessments have collapsed into a cheating arms race. The diffusion of capable coding agents such as Cursor, Claude Code, and an entire shadow ecosystem of "interview-coder" overlays designed to be undetectable to proctors, has made it trivial for a moderately resourceful candidate to pass a take-home or live-coding screen without exercising the underlying skill the screen is supposed to measure. LeetCode-style assessments now produce near-uniformly high scores from candidates whose unaided performance varies wildly. Take-home projects, once a defensible work-sample, have become indistinguishable from agent output. The signal-to-noise ratio of the entire screening layer has degraded to the point where many engineering leaders treat coding-test scores as evidence of nothing but the candidate's willingness to attempt the test. The instrument has not just lost calibration; it has lost the ability to discriminate at all.

The capability that *does* matter, such as engineering effectiveness in an AI-augmented workflow, is not measured anywhere. The day-to-day work of a 2026 software engineer is a continuous negotiation with AI tooling: deciding when to delegate, how to scope a prompt, how to verify generated output, when to override an agent's plan, when to abandon agent-assisted work and write code by hand. Engineers who do this well dramatically outperform those who do not, but no current interview instrument generates a defensible signal about this capability. Banning AI from the interview selects for unaided coding while leaving the AI-orchestration calibration skill entirely untested. Hiring teams therefore ship engineers into AI-native organizations without knowing whether those engineers can actually function in them.

These are the proximate failures. They sit on top of an older, structural failure: hiring instruments invented a century ago that select for the misleading proxy of *interview preparedness* and are structurally blind to the qualities that drive performance in conditions of real uncertainty. The combination is acute. Hiring processes optimized for the wrong signals produce systematic mismatches between role requirements and the people placed into them. These are invisible at the point of hire, but manifest downstream in slower execution, higher management overhead, and accelerating attrition.

For corporations, the consequence is that evaluation frameworks cannot reliably identify who will drive product velocity, navigate ambiguity, or lead effective human–AI collaboration; organizations therefore lose their ability to build the capability portfolios that competitive advantage depends on. For society, as AI increasingly handles execution-layer tasks, the skills that generate human economic value are shifting toward judgment, synthesis, ethical reasoning, and collaborative intelligence. Yet talent pipelines continue to screen for the skills AI is replacing fastest, while filtering out the human qualities that remain irreplaceable. Over time, this produces a workforce-allocation problem of significant scale.

We are therefore in urgent need of a critical reassessment of technical talent evaluation, and an innovative reinvention of the technical interview as an instrument.

2.2 Product Scope: The Technical Layer Only

Basanite is, deliberately, not attempting to replace the entire interview. Modern hiring funnels generally consist of three evaluative layers:

1. **The technical layer**: does the candidate have the cognitive, technical, and AI-collaboration capabilities the role requires?
2. **The psychometric layer**: personality structure, cognitive style, motivational profile.
3. **The culture-fit / values-alignment layer**: does the candidate's working style match the team's, and do their values align with the organization's mission?

At the current stage, Basanite addresses **only the technical layer**, and it does so end-to-end. Psychometric assessment and culture-fit evaluation are explicitly out of scope for the current product. They are not unimportant. They are different problems with different evidentiary bases, different ethical considerations, different regulatory surfaces, and different commercial dynamics. Bundling them into a single product would dilute the rigor of each.

This is a scope choice rather than a permanent exclusion. Once the technical layer is solved at production quality and Basanite has accumulated the validation data to defend its claims, an extension into psychometric or culture-fit assessment may be revisited as a separate product line.

---

3. Target Domain: The Tech-Industry Map

Basanite's first market is the technical interview *within the tech industry*. Since "Tech industry" is a coarse term; we use it here as shorthand for the set of verticals in which engineering, data, and AI-systems work is the core value-generating activity, or in which, even where technology is not core, the technical population is large enough to sustain a dedicated assessment infrastructure.

The following table enumerates the verticals Basanite targets, the roles within each that constitute Basanite's primary surface area, and the role-shape at junior, mid, and senior level. Every cell is a distinct calibration profile in the product: a weighting across the eight metacognitive dimensions, a depth register (applied vs. research/architecture), and a sandbox library of scenarios appropriate to that role-seniority pair.

3.1 Verticals, Roles, and Seniority Bands

Vertical	Representative Role	Junior (≈ L3 / 0–3 yrs)	Mid (≈ L4–L5 / 3–7 yrs)	Senior (≈ L6+ / 7+ yrs)
Consumer Internet & SaaS	Backend Engineer	Owns small features under guidance; writes service code against an existing framework; reads more code than they write.	Owns end-to-end services; makes design tradeoffs; mentors juniors; influences team-level technical decisions.	Sets architectural direction across services; owns reliability and scaling outcomes; sets technical hiring bar.
	Full-Stack Engineer	Implements UI components and CRUD endpoints.	Owns features across the stack; debugs across	Defines product–engineering interface; leads cross-

Vertical	Representative Role	Junior (≈ L3 / 0–3 yrs)	Mid (≈ L4–L5 / 3–7 yrs)	Senior (≈ L6+ / 7+ yrs)
			boundaries; designs API contracts.	functional initiatives; sets stack-level technical strategy.
	Frontend Engineer	Builds isolated components; consumes APIs.	Owns app-level state and performance; sets component standards; leads design-engineering collaboration.	Defines design-system architecture; leads platform-wide rendering strategy; influences product direction through technical possibility.
	Mobile Engineer	Implements screens against existing patterns.	Owns features end-to-end; understands platform internals; leads release engineering.	Sets mobile architecture; influences platform-team dependencies; owns native–cross-platform strategy.
Cloud Infrastructure & DevOps	Site Reliability Engineer	Responds to alerts; runs runbooks; writes basic automation.	Owns service-level objectives; designs observability; leads incident response.	Designs platform-wide reliability; sets SLO/error-budget policy; leads org-wide incident programs.
	Platform Engineer	Maintains internal tooling; supports developer workflows.	Designs CI/CD, IaC, and developer-experience platforms.	Owns multi-cluster, multi-region platform strategy; defines internal-developer-platform vision.
	Cloud Solutions Engineer	Configures cloud resources to specification.	Designs cloud architectures for product workloads; owns cost and reliability.	Sets cloud strategy across the organization; leads multi-cloud and migration programs.
AI / ML / Data	ML Engineer	Trains models on prepared data; deploys via existing pipelines.	Owns training and serving pipelines; designs evaluation; debugs production model behaviour.	Defines modeling strategy; owns the research-to-production pathway; leads modeling roadmap.

Vertical	Representative Role	Junior (≈ L3 / 0–3 yrs)	Mid (≈ L4–L5 / 3–7 yrs)	Senior (≈ L6+ / 7+ yrs)
	Applied AI Engineer	Wires LLM APIs into product features; iterates on prompts.	Designs agentic systems, retrieval architectures, and evaluation harnesses.	Sets organization-wide AI-product strategy; designs novel agent architectures; influences foundational tooling choices.
	Research Engineer	Implements ideas from papers under guidance.	Co-leads research projects; bridges research and engineering.	Sets research agenda; owns publication and IP strategy.
	Data Engineer	Builds ETL jobs against existing schemas.	Owns data models, warehouse design, pipeline reliability.	Defines data-platform architecture; sets governance and lineage standards.
	Data Scientist / Analyst	Runs analyses against existing datasets; builds dashboards.	Owns experimental design; builds production models; advises product.	Leads research agenda; influences product strategy with quantitative work.
Cybersecurity	Security Engineer	Triages alerts; runs scans; writes detections.	Designs detection coverage; leads incident response; threat-models systems.	Sets security architecture; owns red-team / blue-team strategy.
	Application Security Engineer	Reviews code for common vulnerability classes.	Owns SDLC security; designs sec-tooling; leads remediation programs.	Sets product-security strategy across engineering org.
	Offensive Security / Red Team	Executes assessments under guidance.	Leads engagements; develops custom tooling and tradecraft.	Sets adversary-emulation strategy; advises board-level risk.
Fintech & Financial	Backend Engineer (Payments / Risk)	Implements transactional logic against existing rails.	Owns ledger correctness, idempotency,	Sets payments / risk architecture; navigates

Vertical	Representative Role	Junior (≈ L3 / 0–3 yrs)	Mid (≈ L4–L5 / 3–7 yrs)	Senior (≈ L6+ / 7+ yrs)
Services Tech			reconciliation, and audit surfaces.	regulatory surface.
	Quantitative Developer	Implements models from research notebooks.	Owns research-to-production pipeline; optimizes execution.	Defines firm-level quant infrastructure; leads modeling-engineering interface.
	Trading Systems Engineer	Maintains low-latency components.	Owns subsystems of the trading stack; tunes performance.	Sets trading-platform architecture; owns latency and correctness contracts firm-wide.
HealthTech & BioTech	Clinical Software Engineer	Builds against existing EHR/HL7/FHIR integrations.	Owns regulated software components (SaMD, IEC 62304).	Sets architecture for regulated systems; navigates FDA/MHRA pathways.
	Bioinformatics Engineer	Runs established pipelines; QCs outputs.	Designs novel analysis pipelines; integrates wet-lab and dry-lab.	Sets computational-biology platform strategy; leads cross-modality programs.
Hardware & Semiconductors	Embedded / Firmware Engineer	Implements drivers under supervision.	Owns firmware modules; debugs at the hardware-software boundary.	Defines firmware architecture; co-designs with silicon teams.
	Silicon / Verification Engineer	Writes testbenches against specs.	Owns block-level verification; designs coverage.	Sets verification methodology; owns sign-off.
Robotics & Autonomous Systems	Robotics Software Engineer	Implements perception/control modules from designs.	Owns subsystem (perception, planning, control).	Sets autonomy stack architecture; owns safety case.
	Controls Engineer	Tunes existing controllers.	Designs controllers for new platforms.	Sets controls-architecture and safety-stack strategy.

Vertical	Representative Role	Junior (≈ L3 / 0–3 yrs)	Mid (≈ L4–L5 / 3–7 yrs)	Senior (≈ L6+ / 7+ yrs)
Gaming & Interactive Media	Gameplay / Engine Engineer	Implements gameplay features against existing engine.	Owns engine subsystems (rendering, physics, networking).	Defines engine architecture; leads cross-title platform strategy.
	Graphics / Rendering Engineer	Implements shaders and effects to spec.	Owns rendering pipeline subsystems.	Sets rendering architecture; influences hardware-vendor relationships.
Developer Tools & Languages	Compiler / Language Engineer	Implements optimizations under guidance.	Owns compiler subsystems; designs language features.	Sets language / toolchain strategy.
	Database / Storage Engineer	Implements storage components to spec.	Owns storage subsystems (query, indexing, replication).	Sets database architecture; leads protocol/format design.

3.2 Notes on the Map

- Seniority bands deliberately use a compressed three-level structure (junior / mid / senior) rather than the full L3–L8 ladder used inside large tech employers. Sub-band calibration (e.g., distinguishing L4 from L5, or distinguishing Staff from Senior Staff) is part of the human final-round responsibility, not Basanite's; Section 5.1.4 ("Honest About What AI Can and Cannot Do") explains why.
- Each cell anchors a calibration profile in the product. The same role at different seniority levels weights the eight dimensions differently. For example, *Tacit-Knowledge Articulation* is typically weighted lightly at the junior band and heavily at the senior band, while *Verification Rigor in AI Collaboration* is weighted heavily at every band.
- The map is a living artifact. New verticals (e.g., quantum software, BCI engineering) and new roles (e.g., AI evaluation engineer, prompt-systems engineer, agent-platform engineer) will be added as their job markets reach the volume threshold at which dedicated calibration is justified.
- Roles outside this map such as sales, marketing, operations, legal, and executive, are **out of scope** for V2.0 (see Section 5.2 in the V1.0 customer-targeting framework, preserved below in Section 9.2).

4. Human Intelligence That Matters: The Eight Metacognitive Dimensions

Research spanning cognitive science, philosophy of knowledge, behavioral decision theory, organizational psychology, and the emerging literature on human–AI collaboration converges on a different account of what drives genuine value in technical organizations. The qualities that distinguish high performers in complex, adaptive, AI-era work environments are not the ones that conventional technical-interview instruments capture. Basanite operationalizes eight such dimensions. Each is given here with a formal definition, the construct's intellectual provenance, and a pointer to the empirical or theoretical literature it rests on.

These eight dimensions share a structural feature that explains why conventional hiring instruments cannot detect them: they cannot be retrieved from a knowledge base. They are not facts to be recalled or algorithms to be executed. They are capabilities forged through real experience, expressed through real judgment, and legible only to evaluators who know what to look for and how to ask.

4.1 Judgment Under Ambiguity

Definition. The capacity to commit to a defensible course of action when the available information is incomplete, the stakes are non-trivial, and neither paralysis nor false confidence is acceptable. Operationally: the candidate distinguishes irreducible uncertainty from resolvable uncertainty, allocates effort to the latter without fetishizing the former, and acts.

Intellectual provenance. The construct sits at the intersection of Frank Knight's foundational distinction between *risk* (where probabilities are knowable) and *uncertainty* (where they are not), the literature on the calibration of expert judgment, and naturalistic decision theory.

Key references. Knight, F. H. (1921). *Risk, Uncertainty, and Profit*. Tetlock, P. E. (2005). *Expert Political Judgment: How Good Is It? How Can We Know?* Princeton University Press. Tetlock, P. E., & Gardner, D. (2015). *Superforecasting: The Art and Science of Prediction*. Crown.

Why it matters now. AI tools have absorbed a large share of routine analysis. The residual human contribution is increasingly the call made when analysis bottoms out when no further inquiry will sharpen the answer and a decision must nonetheless be taken.

4.2 Tacit-Knowledge Articulation

Definition. The ability to surface, interrogate, and partially transmit knowledge that lives in practice rather than in text.

Intellectual provenance. Polanyi's canonical claim that "we can know more than we can tell" is the origin point. Nonaka and Takeuchi extended the construct into organizational knowledge creation; Collins gave it a more analytic taxonomy.

Key references. Polanyi, M. (1958). *Personal Knowledge: Towards a Post-Critical Philosophy*. University of Chicago Press. Polanyi, M. (1966). *The Tacit Dimension*. Doubleday. Nonaka, I., & Takeuchi, H. (1995). *The Knowledge-Creating Company*. Oxford University Press. Collins, H. (2010). *Tacit and Explicit Knowledge*. University of Chicago Press.

Why it matters now. A senior engineer's leverage in an AI-native organization increasingly comes from the ability to convert tacit pattern-recognition into prompts, evaluation rubrics, and review heuristics that scale across the team.

4.3 Intuition Under Data Scarcity

Definition. The capacity for rapid, accurate situational assessment in conditions where deliberative analysis is infeasible. The recognition-primed judgment that distinguishes genuine experts from those who merely know the vocabulary of expertise.

Intellectual provenance. Klein's recognition-primed decision (RPD) model, developed across two decades of fieldwork with firefighters, military commanders, and ICU clinicians, is the foundational source. The Dreyfus brothers' five-stage skill-acquisition model gives the developmental account. The integration with Kahneman's earlier skepticism toward expert intuition is captured in their joint paper.

Key references. Klein, G. (1998). *Sources of Power: How People Make Decisions*. MIT Press. Klein, G. (2008). "Naturalistic Decision Making." *Human Factors*, 50(3), 456–460. Kahneman, D., & Klein, G. (2009). "Conditions for Intuitive Expertise: A Failure to Disagree." *American Psychologist*, 64(6), 515–526. Dreyfus, H. L., & Dreyfus, S. E. (1986). *Mind Over Machine*. Free Press.

Why it matters now. Production incidents, critical design choices, and on-call decisions are still made under time pressure that rules out exhaustive analysis. The engineer whose intuition is calibrated to the system at hand is the one who keeps the system up.

4.4 Psychological Safety and Collective Learning

Definition. The interpersonal stance (and the team-level conditions one creates around oneself) under which errors surface early, dissent is voiced, and the team's collective intelligence exceeds the sum of its individuals.

Intellectual provenance. Edmondson's seminal study of hospital teams established the construct empirically; her later work scaled it into a theory of organizational learning. Google's Project Aristotle later corroborated the finding in software-engineering teams at scale.

Key references. Edmondson, A. (1999). "Psychological Safety and Learning Behavior in Work Teams." *Administrative Science Quarterly*, 44(2), 350–383. Edmondson, A. (2018). *The Fearless Organization: Creating Psychological Safety in the Workplace for Learning, Innovation, and Growth*. Wiley. Duhigg, C. (2016). "What Google Learned From Its Quest to Build the Perfect Team." *The New York Times Magazine*, February 25, 2016.

Why it matters now. Effective human–AI teams require *more* error-surfacing than human-only teams: AI outputs are confidently wrong in ways human outputs rarely are, and the team's only defense is a culture that catches them before they compound.

4.5 Creative Problem Reframing

Definition. The ability to recognize when the team is solving the wrong problem, and to articulate a reformulation that opens new solution space. This is a capacity that is actively suppressed by evaluation systems that reward convergence on expected answers.

Intellectual provenance. Schön's distinction between *problem-setting* and *problem-solving* identified reframing as the irreducible expert contribution. Dorst's frame-innovation methodology operationalized the practice for design contexts. Csikszentmihalyi's program on the psychology of creativity provides the broader empirical context.

Key references. Schön, D. A. (1983). *The Reflective Practitioner: How Professionals Think in Action*. Basic Books. Dorst, K. (2015). *Frame Innovation: Create New Thinking by Design*. MIT Press. Csikszentmihalyi, M. (1996). *Creativity: Flow and the Psychology of Discovery and Invention*. Harper Perennial.

Why it matters now. Most of the value created in AI-augmented engineering comes from asking a better question of the agent, not from extracting a better answer to the existing one.

4.6 Ethical Reasoning in Practice

Definition. The capacity to feel the weight of real tradeoffs to a user, a colleague, or a downstream system and to navigate them with integrity when the right answer is costly.

Intellectual provenance. Aristotle's *phronesis* (practical wisdom) is the philosophical antecedent. Rest's Four-Component Model and Defining Issues Test produced the dominant 20th-century empirical research program on moral reasoning; Haidt's social-intuitionist model offered an influential rival account. Recent applied work on AI ethics including its limits as a deliverable artifact, and its embedding in engineering practice, extends the literature into our specific domain.

Key references. Rest, J. R. (1986). *Moral Development: Advances in Research and Theory*. Praeger. Haidt, J. (2001). "The Emotional Dog and Its Rational Tail: A Social Intuitionist Approach to Moral Judgment." *Psychological Review*, 108(4), 814–834. Mittelstadt, B. D., Allo, P., Taddeo, M., Wachter, S., & Floridi, L. (2016). "The Ethics of Algorithms: Mapping the Debate." *Big Data & Society*, 3(2). Metcalf, J., Moss, E., & boyd, d. (2019). "Owning Ethics: Corporate Logics, Silicon Valley, and the Institutionalization of Ethics." *Social Research*, 86(2), 449–476.

Why it matters now. The artifacts engineers ship have ethical surfaces that earlier generations of software did not. The capacity to feel a tradeoff is the precondition for navigating it.

4.7 Transformative Learning From Experience

Definition. The metacognitive self-awareness that separates those who *accumulate* experience from those who *learn* from it: the ability to update prior beliefs in response to disconfirming evidence, and to do so in proportion to the evidence's weight.

Intellectual provenance. Flavell introduced metacognition as a research construct in cognitive psychology. Kolb's experiential-learning cycle is the canonical developmental model. Mezirow's theory of transformative learning addresses the deeper case of belief restructuring. Argyris and Schön's *double-loop learning* maps the construct onto organizational practice.

Key references. Flavell, J. H. (1979). "Metacognition and Cognitive Monitoring: A New Area of Cognitive-Developmental Inquiry." *American Psychologist*, 34(10), 906–911. Kolb, D. A. (1984). *Experiential Learning: Experience as the Source of Learning and Development*. Prentice-Hall. Mezirow, J. (1991). *Transformative Dimensions of Adult Learning*. Jossey-Bass. Argyris, C., & Schön, D. A. (1978). *Organizational Learning: A Theory of Action Perspective*. Addison-Wesley.

Why it matters now. The half-life of technical knowledge in 2026 is shorter than at any prior moment in the industry's history. The engineer who cannot restructure their working assumptions on a 12-month cycle is functionally obsolete.

4.8 Human–AI Collaboration Intelligence

Definition. The skill of producing high-quality outcomes through fluent, judicious orchestration of AI tooling: knowing when to delegate, how to scope and decompose tasks for an agent, how to evaluate and verify agent output, when to override the agent, and when to abandon agent-assisted work in favor of unaided execution. It is *calibration* that constitutes the skill.

Intellectual provenance. Daugherty and Wilson provided the early industrial taxonomy of human–AI collaboration patterns. Brynjolfsson framed the strategic question of whether AI is deployed to substitute for or augment human labor. Mollick's work on "co-intelligence" and the empirical study by Dell'Acqua and colleagues on the "jagged technological frontier" provide the most current framing, including the now-canonical finding that AI tools improve performance on tasks within the frontier and *degrade* performance on tasks outside it. The asymmetry makes calibration itself the high-leverage variable.

Key references. Daugherty, P. R., & Wilson, H. J. (2018). *Human + Machine: Reimagining Work in the Age of AI*. Harvard Business Review Press. Brynjolfsson, E. (2022). "The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence." *Daedalus*, 151(2), 272–287. Dell'Acqua, F., McFowland, E., Mollick, E. R., Lifshitz-Assaf, H., Kellogg, K., Rajendran, S., Krayer, L., Candelon, F., & Lakhani, K. R. (2023). "Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of AI on Knowledge Worker Productivity and Quality." Harvard Business School Working Paper No. 24-013. Mollick, E. (2024). *Co-Intelligence: Living and Working with AI*. Portfolio.

Why it matters now. This is the dimension that no existing technical-interview instrument measures, and it is the one most predictive of on-the-job performance in 2026. Basanite's Round 2 (Section 7.3) is purpose-built to generate signal on it.

---

5. Design Philosophy

Assess Genuine Capability, Not Performed Competence

The central conviction behind every design decision at Basanite is simple: most hiring tools are optimised for the wrong signal. They measure how well someone can approximate the idea of a good candidate, rather than whether they actually are one.

5.1 Five Principles

5.1.1 Depth Over Breadth

Basanite is designed around progressive excavation: each layer of the assessment exists to move one level deeper into signal quality. A candidate who answers fluently at the surface should encounter ground that shifts beneath them at the next layer. The edges of real ability are blurry. Performed ability has no edges.

5.1.2 Structure as a Fairness Mechanism

Unstructured evaluation produces advantage for those who are most culturally legible to the evaluator. Standardized evalution encourages cheating and rewards those who perform. In the absence of valid capability signals, assessors default to cultural familiarity. Candidates who mirror the background, communication style, and social register of the interviewer are advantaged not because they are more capable, but because they are easier to read.

Basanite's structure aim to be equalising. By anchoring every evaluation to consistent frameworks, principled questioning logic, and explicit scoring criteria grounded in observable evidence, the system creates conditions where a self-taught engineer without institutional pedigree can be seen as clearly as one from a target university, as both are being asked the same questions in the same spirit, with the same depth of follow-up.

5.1.3 The Encounter as Part of the Assessment

The experience a candidate has inside a Basanite evaluation should feel like a rigorous but genuinely curious conversation instead of an interrogation, a performance, or a test of how well they have prepared for exams. This matters for two reasons:

1. The collaborative, uncertainty-tolerant environments that the best technical organisations are building require people who can think under genuine pressure, not people who can perform competence in artificial conditions. An evaluation system that rewards the latter is selecting against the former.
2. Candidates who feel genuinely seen regardless of outcome become the most powerful distribution channel available. The reputation of an evaluation process travels through the talent communities it touches. A process that candidates respect is one that attracts the candidates worth reaching.

5.1.4 Honest About What AI Can and Cannot Do

Basanite will never claim more than the evidence supports. The system is strongest where AI has a genuine structural advantage over human evaluation: in consistency, in scale, in the ability to follow a thread across a long conversation without fatigue or drift, and in the capacity to maintain identical standards regardless of the order in which candidates are seen. These are real advantages, and they compound. The tenth candidate in a Basanite evaluation receives the same quality of assessment as the first. This is not true of human-led processes, and the difference matters at scale.

But there are things that require irreducible human judgment, such as nuances of seniority assessment, domain-specific technical verification, and the final hiring decision itself. Basanite is designed to be explicit about these limits: flagging where human expertise is required, producing quotable evidence rather than opaque scores, and positioning itself as the infrastructure that makes human judgment better rather than the mechanism that replaces it. Intellectual honesty about capability boundaries is a design requirement, not a marketing disclaimer.

5.1.5 Evaluation as a Two-Way Mirror

The best hiring processes leave candidates with a clearer understanding of themselves instead of just a verdict. A process that produces genuine signal about a candidate's capabilities should, by definition, produce signal that is legible to the candidate as well as the organisation. If a candidate cannot understand what was being assessed and why, the assessment itself is suspect.

This principle has a practical implication for system design: every evaluation strategy deployed by Basanite must be one that can be honestly explained to the candidate it is applied to. Strategies that depend on deception to function that work only because the candidate does not know what is being measured, are excluded not merely on ethical grounds, but on epistemic ones. An evaluation that cannot withstand transparency is not producing the signal it claims to produce.

---

6. Prototype Architecture

[Preserved from V1.0, content to be filled in.]

---

7. Interview Methodology: A Two-Round Architecture

A central change in V2.0 is the move from a single-round assessment to a two-round architecture. Each round generates a class of signal the other cannot. The composite is what produces a defensible technical-hire recommendation.

7.1 Why Two Rounds

The V1.0 design relied entirely on a structured conversational assessment. That instrument is strong on dimensions that surface through narrative and reasoning under follow-up. It is weak on the dimension that has become most economically important: how the candidate actually performs when they sit down at a keyboard, with a coding agent and a real codebase, and have to ship.

Conversation reveals what a candidate *thinks*. Observed work reveals what they *do*. The gap between the two is itself diagnostic.

The two rounds are deliberately complementary:

- **Round 1: Structured Conversational Assessment.** Generates signal across the cognitive, judgmental, and tacit-knowledge dimensions (4.1, 4.2, 4.3, 4.5, 4.6, 4.7).
- **Round 2: AI Collaboration Workbench.** Generates signal across executional, AI-orchestration, verification-rigor, and engineering-taste dimensions (4.4, 4.7, 4.8, and triangulates 4.1 and 4.5 against Round 1).

A candidate's composite hirer report integrates both rounds. The two are not redundant; they triangulate.

7.2 Round 1: Structured Conversational Assessment

This is the assessment described in V1.0 and preserved without substantive change. Section 11 (Technical Innovation and Differentiation) catalogs the techniques the conversational round employs in detail: narrative anchoring, experience-grounded questioning, open sandbox questioning, decision tracing, boundary-condition probing, counterfactual pressure, forced commitment, compulsory tradeoff framing, progressive excavation, vagueness targeting, fourth-wall break, honest failure elicitation, predict-your-own-error, regret probing, narrative consistency tracking, parameter verification, cognitive priority testing, information-gap injection, structural opacity, AI-output signal detection, cognitive-load escalation, latency awareness, and tacit-knowledge consistency testing.

The conversational round runs typically 20-30 minutes depending on role seniority and budget tier. It terminates on signal saturation, not on question count.

7.3 Round 2: The AI Collaboration Workbench

Round 2 is new in V2.0. It places the candidate in a sandboxed environment that closely resembles their target job, hands them an AI coding agent and a realistic codebase, and observes how they ship. The design pioneeres the practice of testing engineers *with* AI rather than *against* it.

7.3.1 Form Factor

The candidate is provisioned with:

- **A role-matched codebase.** The codebase is not a toy. It is a multi-thousand-line synthetic project calibrated to the target vertical and seniority level (drawn from the role row in Section 3). A candidate interviewing for a backend SaaS role works against a SaaS codebase with realistic service boundaries, schema migrations, and feature-flag plumbing. A candidate interviewing for an applied-AI role works against an agentic-systems codebase with retrieval, tool-use, and evaluation harnesses already in place. A candidate interviewing for a security-engineering role works against a codebase with seeded vulnerabilities and a partial detection stack.
- **A real ticket or a small bundle of tickets.** Tickets are written in the style and granularity the candidate would receive on day one of the role. Granularity is calibrated to seniority: junior tickets are bounded and well-specified; mid-level tickets contain a non-trivial design choice; senior tickets are deliberately under-specified and require the candidate to scope, sequence, and (in some scenarios) negotiate with a simulated requester before doing any code work.
- **An AI coding agent of their choice.** The candidate may use Claude Code, Cursor, Copilot, Aider, or a local CLI agent. Basanite is tooling-agnostic. The candidate works the way they actually work, not in a synthetic chat interface. (This is a deliberate departure from products that constrain candidates to a custom UI: artificial constraints distort the signal.)
- **A clean execution environment.** A web-based VS Code instance with a terminal, git, the codebase preloaded, and the ability to run tests and build commands. The environment is sandboxed, instrumented, and otherwise transparent.

7.3.2 Duration and Pacing

The session is time-boxed, though the duration can be further customized for each end user, the typical choices are:

- 35 minutes for junior-band roles.
- 60 minutes for mid-band roles.
- 90 minutes for senior-band roles, with optional extension to 120 minutes for architecture-heavy variants.

A 5-minute orientation precedes the timer. A 10-minute structured reflection conversation follows (7.3.4). Total seat time: roughly 75–135 minutes depending on band.

7.3.3 Instrumentation

Throughout the session, Basanite captures:

- **Full keystroke and command history** in the terminal and editor.
- **The complete dialogue between the candidate and their AI agent**: prompts issued, responses received, suggestions accepted, suggestions rejected, suggestions reverted.
- **Git state at fine-grained intervals**: every commit, plus periodic snapshots between commits.
- **Time-on-task per file and per ticket subtask.**
- **Verification behavior**: did the candidate run tests, write new tests, exercise the change manually, or none of the above before considering the work complete?

This instrumentation is disclosed to the candidate at the start of the session, in plain language. The structural-opacity principle (5.1.2 / 5.1.5) governs *which dimensions are scored from the trace*; the existence of the trace itself is transparent. This is consistent with the V1.0 commitment that no Basanite assessment depends on deception to function.

7.3.4 The Reflection Conversation

After the timed session, Basanite's interview agent re-engages the candidate in a 10-minute structured conversation. The conversation has a fixed shape:

1. *"Walk me through what you did and why."* The candidate narrates their own session. The narration is cross-referenced against the trace.
2. *"Where did you and the agent disagree, and how did you resolve it?"* This surfaces the candidate's calibration of when to trust the agent and when to override it.
3. *"What would you have done differently with another 30 minutes?"* This exposes engineering taste: what the candidate sees as incomplete in their own work.
4. *"Walk me through one specific moment where the agent produced output you accepted. Why did you trust it there?"* The answer is cross-checked against the trace at the timestamp the candidate identifies.

The reflection conversation closes the gap between *what the candidate did* (observable in the trace) and *what they understood themselves to be doing* (only available through narration). Discrepancies between the two are themselves signal, and frequently the most diagnostic signal Round 2 produces.

7.3.5 Scoring Dimensions Specific to Round 2

Six sub-dimensions are scored from the Round 2 trace and reflection. The first four sit primarily within the Human–AI Collaboration Intelligence dimension (4.8); the last two probe engineering capability that conversational assessment cannot reach.

Sub-dimension	What is being scored	Primary evidence
Delegation calibration	Does the candidate route work to the agent where the agent has comparative advantage, and do work themselves where they do?	Prompt-log content vs. keystroke ratio per ticket subtask
Prompt quality and decomposition	Are prompts well-scoped, contextually situated, and decomposed appropriately? Or are they large, vague, and one-shot?	Prompt-log analysis; iteration patterns
Verification rigor	Does the candidate verify generated output before integrating it, and through what means (reading, tests, manual exercise)?	Test runs, file diffs, time spent reading vs. writing
Override judgment	When the agent is wrong, does the candidate notice? When does the candidate abandon agent-assisted work?	Reverts, manual rewrites of agent-produced code, fallback patterns
Engineering taste	Is the final solution complete, idiomatic, and consistent with the codebase's conventions?	Final code state vs. codebase priors
Solution completeness	Did the candidate finish the ticket as specified, or did they declare done prematurely?	Ticket acceptance criteria vs. final state

Each sub-dimension is scored with the same quotation-anchored discipline as Round 1: high scores require specific, citable evidence from the trace or the reflection conversation.

7.3.6 What Round 2 Deliberately Does *Not* Measure

Round 2 is not an algorithmic-puzzle test in disguise. The codebase contains no LeetCode-style problems. Tickets are routine engineering tasks, the kind of work the candidate would do every day in the role. The point is not whether the candidate can solve a hard, isolated problem under artificial constraints; it is whether the candidate can ship calibrated, complete work in the way the actual job is done.

Round 2 also does not test whether the candidate uses AI "more" or "less." The target is *judicious* use, calibrated to where the agent helps and where it does not.

7.3.7 Anti-Cheating Posture

Round 2 deliberately inverts the standard anti-cheating posture. Rather than preventing the candidate from using AI, Basanite *requires* it and instruments it. This eliminates the entire "did the candidate use AI" cheating vector that has corrupted contemporary coding screens. The cheating risks that remain are different in shape, such as a third party operating the candidate's machine remotely, or a candidate having someone else complete the session. They are addressed through:

- Identity verification at session start.
- Behavioral biometrics (typing cadence, mouse movement signatures) sampled across the session and compared against the candidate's Round 1 baseline.
- A randomized in-session check-in protocol where the candidate is asked, mid-session, to explain a specific decision they just made. Genuine candidates explain fluently from working memory; substituted operators do not.

7.3.8 Composite Reporting Across Rounds

The hirer report integrates Round 1 and Round 2 into a single document. Where the two rounds agree, the signal is reinforced. Where they disagree, for example, a candidate who articulates strong principles in Round 1 but ships sloppily in Round 2, or vice versa, the disagreement is itself surfaced as a flagged signal for the human final-round interviewer to probe. The composite report explicitly identifies these cross-round discrepancies and recommends interview directions to resolve them.

---

8. (V1.0 Section 7) Technical Innovation and Differentiation

> The following technique inventory describes the methods used inside **Round 1** of the V2.0 architecture. Round 2 techniques are described in Section 7.3.

8.1 Basanite Interview Technique Inventory

8.1.1 Structural Techniques

Macro Three-Part Structure. Every interview follows a fixed high-level arc: narrative foundation → dimension assessment → stress test. Candidates must establish their own story before probing begins.

Signal Saturation Pacing. Transitions between dimensions are governed by internal confidence thresholds, not question counts. The interviewer moves on when it has enough, not when a quota is met.

Experience Path Bifurcation. The interview forks at the outset based on whether the candidate has relevant prior experience. Path A deploys experience-anchored questions; Path B deploys capability-inference questions with a weighted composite scoring model.

Technical Depth Calibration. Before the interview begins, the agent calibrates to one of two depth registers: applied or research/architecture, based on the job description, which governs the intensity of technical probing throughout.

8.1.2 Questioning Techniques

Narrative Anchoring. The opening phase invites the candidate to describe a recent piece of work in their own words, establishing a shared factual baseline from which all subsequent questions grow organically.

Experience-Grounded Questioning. Questions are tied to the candidate's own stated history rather than hypothetical scenarios. The container is always the candidate's real past; the question only works if they actually did the thing.

Open Sandbox Questioning. Where real experience is absent or insufficient for a particular evaluation metrics, the candidate is placed in a realistic hypothetical scenario. Five specialist sandbox design principles govern these questions to prevent them from being answerable by AI lookup alone.

Decision Tracing. The candidate is asked not only what they chose but why they did not choose the alternatives. The quality of their dismissal of Option B reveals more than their justification of Option A.

Boundary Condition Probing. Rather than asking whether a solution works, the interviewer asks under what conditions it fails. Valid technical knowledge has edges; performed knowledge does not.

Counterfactual Pressure. In the stress test phase, the candidate is asked what would have happened if they had made the opposite decision. This requires sustained reasoning in a direction for which they have no prepared answer.

Forced Commitment. When a candidate gives a framework-level answer, the interviewer demands a specific, defensible position. This disrupts the AI tendency toward balanced non-answers.

Compulsory Trade-Off Framing. Sandbox questions are structured to make "consider all dimensions" a wrong answer strategy. Resource constraints or mutually exclusive choices are built into the scenario to force prioritisation.

8.1.3 Follow-Up and Depth Techniques

Progressive Excavation. Each layer of follow-up is designed to go one level deeper into the same territory rather than expand to new ground. The goal is vertical depth, not horizontal coverage.

Vagueness Targeting. When a candidate gives an answer that sounds complete, the interviewer identifies the single most ambiguous detail within it and asks the candidate to elaborate specifically on that point.

Fourth Wall Break. Mid-scenario, the interviewer pulls the candidate out of the hypothetical and asks what was genuinely the first thing that came to mind, not the answer they organised, but the raw initial reaction. AI has no first reaction; humans do.

Honest Failure Elicitation. The candidate is asked to name a decision in their work that they did not think through fully but that happened to turn out fine anyway. Clean answers to this question are a high-risk signal. Slightly uncomfortable, specific ones are not.

Predict-Your-Own-Error. After a candidate gives a judgment, they are asked where that judgment is most likely to be wrong, not theoretically, but based on their own known patterns of error. Generic risk lists are a signal; personal blind spots are the target.

Regret Probing. For any experience the candidate describes as a success, the interviewer asks what they would have done differently. Authentic experience always contains regret; performed experience does not.

8.1.4 Consistency and Authenticity Verification Techniques

Narrative Consistency Tracking. All specific details mentioned by the candidate: team sizes, timelines, technical parameters, project names, are internally logged and referenced in later questions to test whether the account holds together across the full conversation.

Parameter Verification. For technical claims, the interviewer asks for concrete numbers: dataset sizes, latency figures, team headcount, timeline specifics. Real memory is approximate and slightly uncertain; fabricated or AI-assisted memory tends to be either implausibly precise or implausibly round.

Cognitive Priority Testing. A key detail is buried in a scenario description, one that an experienced practitioner would immediately identify as the most important signal, but that an inexperienced candidate or AI would treat as background noise. The candidate is assessed not on what they answer but on what they notice.

Information Gap Injection. A critical variable in a scenario is deliberately left ambiguous, one that only real experience can resolve through assumption. A candidate with genuine experience names a specific calibration factor; one without produces a generic framework.

8.1.5 Anti-Cheating Techniques

Structural Opacity. Evaluation dimensions are never disclosed to the candidate. Question purposes are concealed. The candidate cannot reverse-engineer what is being measured and optimise their performance toward it.

AI Output Signal Detection. The interviewer maintains a checklist of linguistic and structural patterns associated with AI-generated responses: comprehensive coverage without prioritisation, clean narrative arcs without friction, self-criticism that reads as a growth story, absence of any moment of genuine uncertainty.

Cognitive Load Escalation. The deepest follow-up is reserved for the areas where the candidate is most confident, not most vulnerable. Defences are lowest where people feel they are winning; that is where authentic signal is easiest to surface.

Latency Awareness. Anomalous response timing: answers that arrive too quickly for genuine reflection or too slowly for normal recall, is flagged as a risk indicator.

Tacit Knowledge Consistency Test. A candidate who claims a skill they cannot teach, but then describes its internal mechanism with unusual precision, has produced a self-contradicting account. The inability to articulate is part of the claim; fluent articulation undermines it.

8.1.6 Scoring and Output Techniques

Quotation-Anchored Scoring. No dimension may receive a score above three without a specific verbatim statement from the candidate cited as evidence. Scores are grounded in what was actually said, not in overall impression.

Cheating Risk as Independent Dimension. Cheating risk is scored and reported separately from capability. A candidate who may have used AI assistance but produced strong responses is treated differently from one who produced weak responses. The two assessments do not collapse into each other.

Dual Report Generation. Two separate outputs are produced from the same assessment: a hirer report designed to brief the human final interviewer, and a candidate report designed to be useful without being reverse-engineerable. The same underlying data serves two different communicative purposes.

Technical Capability Mapping. Rather than a single technical score, the output produces a structured map: areas of demonstrated depth, areas of surface fluency without underlying understanding, specific nodes requiring human technical expert verification, and surface-fluent but substantively suspect domains.

---

9. (V1.0 Section 5) Defining Target Customers

9.1 Primary Target Customers

Tier 1: High-Volume Technical Recruiters: Enterprise and Scale-Up

Any organisation running more than 30 to 40 technical hires per year has a structural headcount cost in screening and first-round assessment that Basanite directly displaces. This includes:

- High-growth technology companies at Series B and beyond with dedicated talent teams
- Technology divisions of large financial institutions, retailers, and professional services firms that hire engineers at volume but are not primarily technology companies
- Management consultancies and technology consultancies running graduate and experienced hire pipelines

For this customer, the primary commercial argument is not quality but headcount. One Basanite deployment can replace the screening and first-round assessment work of one to three full-time recruiters. At £60,000 to £90,000 per recruiter fully loaded, the payback period on a Basanite contract is measured in months, not years. Quality improvement is the reason the legal and ethical case for the product holds; cost displacement is the reason it gets purchased.

Tier 2: Staffing Agencies and Recruitment Process Outsourcing Firms

Staffing agencies and RPO firms face acute margin pressure as the market commoditises. Their differentiation, increasingly, must come from the quality and defensibility of the assessment they provide instead of only from their candidate databases and client relationships. A Basanite-powered agency can offer enterprise clients a materially stronger assessment product at lower per-candidate cost than a human-led screening process, while improving their own gross margin by reducing the recruiter hours per placement.

The incentive misalignment exists for agencies whose business model depends on volume throughput of weak candidates. It does not exist for agencies attempting to move upmarket toward retained search and quality-differentiated assessment services. The latter is a growing segment, and Basanite is a natural infrastructure partner for that strategic shift.

Tier 3: SMEs Without Dedicated HR Functions

Small and medium-sized technology companies, typically between 20 and 150 employees, often have no dedicated HR or recruiting function at all. Hiring is done directly by founders, CTOs, or engineering leads who do not have the bandwidth for structured screening and lack the institutional knowledge to design rigorous first-round assessments. These companies consistently over-rely on personal networks and referrals, which produces the homogeneous teams that the research literature predicts: brittle, culturally narrow, and poorly calibrated for the uncertainty of early-stage growth.

For this customer, Basanite is not replacing existing headcount, it is providing a capability that does not exist. The commercial framing is closer to infrastructure than to labour substitution: Basanite gives a 30-person company the evaluation sophistication of a company ten times its size, at a fraction of the cost of hiring a Head of Talent to build that sophistication internally. This segment is price-sensitive but large, and the switching cost once embedded in a hiring workflow is substantial.

Tier 4: Universities, Graduate Employers, and Professional Training Programmes

As previously identified, this segment provides both revenue and strategic data value. The revised commercial argument here is more pointed: universities and large graduate employers are running assessment centres that cost £500 to £2,000 per candidate in assessor time, venue costs, and coordination overhead. For a graduate scheme processing 2,000 applications down to 50 offers, the assessment infrastructure cost alone can exceed £500,000 per cycle. Basanite can execute the first two to three stages of that funnel at a fraction of the cost, with demonstrably superior consistency.

Tier 5: Geographically Distributed and Remote-First Organisations

A structural shift in how technical work is organised has created a specific customer need that was not present five years ago. Organisations hiring across multiple time zones and jurisdictions face a coordination cost in scheduling and running human-led assessments that compounds with geographic distribution. Asynchronous AI-led evaluation removes this friction entirely. For a company hiring engineers across three continents from a central talent function, the operational value of an always-available, timezone-agnostic assessment system is significant and immediate.

9.2 Who Basanite Is Not For At This Stage

Non-technical roles, at least initially. Basanite's assessment architecture is purpose-built for technical capability evaluation. The question design logic, the tacit knowledge detection framework, the cheating identification signals, and the transfer testing methodology all assume a candidate with a technical background being evaluated for a technical role. Extending to sales, operations, legal, or marketing roles is not impossible, but it would require a substantive rebuild of the core assessment logic, not a surface-level adaptation. This is a deliberate scope constraint, not a permanent exclusion.

Roles where the evaluation is primarily relational or political. Senior executive hiring is governed by dynamics that no structured assessment system currently handles well. These decisions involve stakeholder alignment, organisational politics, and relationship capital that are genuinely resistant to systematic evaluation. Human judgment, executive search relationships, and reference networks remain the dominant signal sources in this tier. Basanite may eventually have a role here, but it is not the right first market.

Organisations that have already decided they want a human face on every candidate interaction. Some companies, particularly those that market heavily on employer brand and candidate experience, have made a deliberate choice that every touchpoint in their hiring process will be human-led. This is a values position, not a capability gap, and it is not one that Basanite should attempt to argue against. The organisations that hold this position most strongly tend to be in consumer-facing industries where brand perception is tightly managed. They are not the right early adopters.

---

10. (V1.0 Section 6) User Journeys

> Stage 2 of the candidate's journey now spans the two interview rounds described in Section 7. The hirer's journey is unchanged in structure; the budget tier selected in Stage 2 of the hirer's journey controls whether Round 2 is activated for the role.

10.1 The Hirer's Journey

Stage 1: Account Setup and Job Configuration

The hirer logs into their Basanite account and navigates to the job creation flow. They paste or compose a job description in the same format they would use on LinkedIn, Greenhouse, or any standard hiring platform: role title, responsibilities, requirements, and any contextual information about the team or company. No additional formatting or translation is required. Basanite accepts the job description as-is.

Basanite analyses the job description and surfaces a recommended set of evaluation dimensions drawn from its assessment library. The recommendation is generated by matching the role's requirements, seniority signals, and technical domain against Basanite's dimension taxonomy. For a senior backend engineer role, for example, the system might recommend dimensions including technical judgment depth, judgment under ambiguity, and the ability to surface tacit knowledge; for a junior data analyst role, it might weight intuition under data scarcity and creative problem reframing differently. The hirer sees the recommended combination with a brief rationale for each dimension.

The hirer can accept the recommendation as a complete package, remove dimensions they consider less relevant to this specific role, or browse the full dimension library and add dimensions not included in the default recommendation. The interface presents dimensions with plain-language descriptions and example signal indicators, so the hirer does not need to understand the underlying assessment methodology to make informed choices.

Stage 2: Eligibility Constraints and Budget

Once the evaluation dimensions are confirmed, Basanite prompts the hirer to specify any hard eligibility constraints that should gate candidates before assessment begins. These include:

- Right to work in a specific jurisdiction
- Graduate or current student status
- Minimum or maximum years of professional experience
- Specific technical prerequisites and how negotiable they are for the role

- Location or timezone requirements

These constraints are stored as pre-screening filters. Candidates who do not meet them are redirected without consuming assessment capacity, and receive a respectful notification that they do not meet the eligibility criteria for this specific role at this time.

Basanite then asks the hirer to specify their assessment budget, either as a total spend for the hiring round or as a per-candidate cap. This budget determines the depth and scope of assessment each candidate receives: the number of evaluation dimensions actively probed, the length of the interview, and whether advanced features such as transfer testing or knowledge reproduction assessment are activated. The system presents the hirer with a clear mapping between budget tiers and assessment scope before they confirm.

Stage 3: Role-Specific Prompt Generation

With job description, evaluation dimensions, eligibility constraints, and budget confirmed, Basanite generates a role-specific interview agent prompt. This prompt is derived from Basanite's master assessment framework and calibrated to the specific parameters the hirer has set: the target dimensions, the technical depth tier appropriate to the role's seniority, and the candidate experience path determined by whether candidates are likely to have directly relevant prior experience.

The generated prompt is stored against the role and activates automatically when candidates begin their assessment. The hirer does not see or interact with the prompt directly, it is operational infrastructure, not a configurable document. The hirer can preview the types of questions and assessment areas the interview will cover in plain language, but the specific prompt logic remains internal to protect assessment integrity.

The role is now live and ready to receive candidates.

10.2 The Candidate's Journey

Stage 1: Application and Onboarding

The candidate arrives at the Basanite assessment link, either directly from a job posting or via an invitation from the hirer. They create a Basanite account or log into an existing one. If they have completed Basanite assessments for other roles previously, their profile history is retained, but each assessment is conducted fresh against the specific role's parameters.

The candidate uploads their CV or résumé. Basanite processes this document and extracts the candidate's experience history, technical background, project work, and any signals relevant to the role being assessed. This information is used to personalise the assessment encounter. The system identifies which of the candidate's experiences are most relevant to the target role and uses these as anchor points for the interview.

Before the assessment begins, Basanite presents the candidate with a brief plain-language explanation of what the assessment involves, approximately how long it will take, and what kinds of questions they can expect. The tone is transparent and respectful: candidates are told they will be asked about their real experiences and their genuine ways of thinking, and that there are no trick questions or expected answers they should be trying to perform toward.

Stage 2: Assessment

> Where the role's budget tier activates the two-round architecture (Section 7), Stage 2 spans both the conversational round and the AI Collaboration Workbench. The candidate completes Round 1 first; Round 2 is scheduled as a separate session, typically within seven days. The narrative below describes Round 1; Round 2's candidate experience is described in 7.3.

The candidate enters the interview. Basanite begins with an open invitation for the candidate to walk through a recent piece of work in their own words: a project, a decision, a problem they solved. This narrative foundation serves multiple functions: it establishes a shared factual basis for follow-up questions, it gives the candidate a comfortable entry point, and it gives the system an initial profile of the candidate's experience depth and communication style against which to calibrate the remainder of the assessment.

From this foundation, Basanite moves through the active evaluation dimensions confirmed for the role. Questions grow organically from the candidate's own narrative rather than from a fixed script. Where the candidate has demonstrated relevant prior experience, the system deploys real-experience questions that probe the depth and authenticity of that experience. Where relevant experience is absent or thin, the system activates open sandbox questions designed to surface the candidate's reasoning and judgment in hypothetical but realistic scenarios.

Throughout the interview, Basanite tracks narrative consistency, monitors for signals that distinguish genuine capability from performance, and calibrates follow-up depth based on the signal confidence it is accumulating in each evaluation dimension. The interview continues until signal confidence thresholds are met across the required dimensions, or until the budget-determined time limit is reached.

The candidate experience throughout is designed to feel like a substantive, curious conversation. Basanite does not rush, does not deploy trick questions, and does not penalise candidates for acknowledging uncertainty. Candidates who say "I don't know" in a way that demonstrates genuine metacognitive awareness are treated differently from candidates who produce confident but hollow answers, because the former is a positive signal instead of a failure.

Stage 3: Dual Feedback Reports

When the assessment concludes, Basanite generates two separate feedback documents.

The hirer report is a structured assessment output containing dimension-by-dimension scores, the specific candidate statements that support each score, flagged areas of particular strength or concern, a cheating risk assessment where applicable, a technical capability map identifying areas of depth and potential blind spots, and a one-paragraph synthesis of what this candidate is likely to be like in a real working environment. This report is designed to function as a briefing document for the hirer's final human-led interview: it tells the interviewer where to probe further, what has already been established, and what remains genuinely uncertain. It does not make a hire or no-hire recommendation, as the hiring decision remains with the human.

The candidate report is a brief, neutral summary of their performance delivered in plain language. It acknowledges what they demonstrated well, identifies areas where they could strengthen their presentation or preparation, and offers constructive suggestions for development. The candidate report is deliberately designed to be informative without being reverse-engineerable: it does not reveal the specific evaluation dimensions used, the scoring criteria, or the question logic underlying the assessment. A candidate cannot use their feedback to game a subsequent Basanite assessment for the same or a different role. The tone is respectful and genuinely useful, a candidate who did not progress should come away with something of value from the process.

Stage 4: Ranking and Queue Management

Once the candidate's report is generated, Basanite's ranking algorithm incorporates the assessment output into a composite score for this candidate against this role. The composite draws on dimension scores, background signals extracted from the CV, stated salary expectations, and any hard eligibility data captured during onboarding.

The candidate is assigned a position in the hirer's live ranking queue for this role. The ranking is dynamic: as new candidates complete assessments, the queue updates automatically. The hirer sees a continuously ordered list of assessed candidates, each with their composite score, a summary of their strongest and weakest dimensions, and a direct link to their full hirer report.

The hirer can sort and filter the queue by individual dimensions, salary band, experience level, or composite score. When they are ready to move candidates forward to a final human-led interview, they select from the queue with the context of Basanite's full assessment in hand, knowing not just where each candidate ranked, but why, and what the human interview should focus on to complete the evaluation.

---

11. Academic Partnerships

[Preserved from V1.0, content to be filled in.]

---

12. Corporation Partnerships

[Preserved from V1.0, content to be filled in.]

---

13. Validation Plans

[Preserved from V1.0, content to be filled in.]

Appendix A. Cross-Reference: V1.0 → V2.0

V1.0 Section	V2.0 Section	Change
1. Executive Summary	1. Executive Summary	Preserved
2.1 The Measurement Gap	2.1 The Failure of the Technical Interview	Rewritten — sharpened around technical-interview-specific failure modes (coding-cheating arms race; unmeasured AI-collaboration ability)

V1.0 Section	V2.0 Section	Change
—	2.2 Product Scope: The Technical Layer Only	New — explicit scoping out of psychometric and culture-fit layers
—	3. Target Domain: The Tech-Industry Map	New section, including verticals × roles × seniority table
2.2 Human Intelligence That Matters	4. The Eight Metacognitive Dimensions	Rewritten with formal definitions, intellectual provenance, and key references; expanded from seven dimensions to eight (added 4.8 Human–AI Collaboration Intelligence)
3. Design Philosophy	5. Design Philosophy	Preserved verbatim
4. Prototype Architecture	6. Prototype Architecture	Preserved (placeholder)
—	7. Interview Methodology: Two-Round Architecture	New — Round 1 (existing conversational round) + Round 2 (AI Collaboration Workbench, openround.ai-informed)
7. Technical Innovation and Differentiation	8. Technical Innovation and Differentiation	Preserved verbatim, with note clarifying that the inventory describes Round 1
5. Defining Target Customers	9. Defining Target Customers	Preserved verbatim
6. User Journeys	10. User Journeys	Preserved with one bridging note in candidate Stage 2 pointing to Section 7
8. Academic Partnerships	11. Academic Partnerships	Preserved (placeholder)
9. Corporation Partnerships	12. Corporation Partnerships	Preserved (placeholder)
10. Validation Plans	13. Validation Plans	Preserved (placeholder)

Appendix B. Reference List

Argyris, C., & Schön, D. A. (1978). *Organizational Learning: A Theory of Action Perspective*. Addison-Wesley.

Brynjolfsson, E. (2022). The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence. *Daedalus*, 151(2), 272–287.

Collins, H. (2010). *Tacit and Explicit Knowledge*. University of Chicago Press.

Csikszentmihalyi, M. (1996). *Creativity: Flow and the Psychology of Discovery and Invention*. Harper Perennial.

Daugherty, P. R., & Wilson, H. J. (2018). *Human + Machine: Reimagining Work in the Age of AI*. Harvard Business Review Press.

Dell'Acqua, F., McFowland, E., Mollick, E. R., Lifshitz-Assaf, H., Kellogg, K., Rajendran, S., Krayer, L., Candelon, F., & Lakhani, K. R. (2023). Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of AI on Knowledge Worker Productivity and Quality. Harvard Business School Working Paper No. 24-013.

Dorst, K. (2015). *Frame Innovation: Create New Thinking by Design*. MIT Press.

Dreyfus, H. L., & Dreyfus, S. E. (1986). *Mind Over Machine: The Power of Human Intuition and Expertise in the Era of the Computer*. Free Press.

Duhigg, C. (2016). What Google Learned From Its Quest to Build the Perfect Team. *The New York Times Magazine*, February 25, 2016.

Edmondson, A. (1999). Psychological Safety and Learning Behavior in Work Teams. *Administrative Science Quarterly*, 44(2), 350–383.

Edmondson, A. (2018). *The Fearless Organization: Creating Psychological Safety in the Workplace for Learning, Innovation, and Growth*. Wiley.

Flavell, J. H. (1979). Metacognition and Cognitive Monitoring: A New Area of Cognitive-Developmental Inquiry. *American Psychologist*, 34(10), 906–911.

Haidt, J. (2001). The Emotional Dog and Its Rational Tail: A Social Intuitionist Approach to Moral Judgment. *Psychological Review*, 108(4), 814–834.

Kahneman, D., & Klein, G. (2009). Conditions for Intuitive Expertise: A Failure to Disagree. *American Psychologist*, 64(6), 515–526.

Klein, G. (1998). *Sources of Power: How People Make Decisions*. MIT Press.

Klein, G. (2008). Naturalistic Decision Making. *Human Factors*, 50(3), 456–460.

Knight, F. H. (1921). *Risk, Uncertainty, and Profit*. Houghton Mifflin.

Kolb, D. A. (1984). *Experiential Learning: Experience as the Source of Learning and Development*. Prentice-Hall.

Metcalf, J., Moss, E., & boyd, d. (2019). Owning Ethics: Corporate Logics, Silicon Valley, and the Institutionalization of Ethics. *Social Research*, 86(2), 449–476.

Mezirow, J. (1991). *Transformative Dimensions of Adult Learning*. Jossey-Bass.

Mittelstadt, B. D., Allo, P., Taddeo, M., Wachter, S., & Floridi, L. (2016). The Ethics of Algorithms: Mapping the Debate. *Big Data & Society*, 3(2).

Mollick, E. (2024). *Co-Intelligence: Living and Working with AI*. Portfolio.

Nonaka, I., & Takeuchi, H. (1995). *The Knowledge-Creating Company: How Japanese Companies Create the Dynamics of Innovation*. Oxford University Press.

Polanyi, M. (1958). *Personal Knowledge: Towards a Post-Critical Philosophy*. University of Chicago Press.

Polanyi, M. (1966). *The Tacit Dimension*. Doubleday.

Rest, J. R. (1986). *Moral Development: Advances in Research and Theory*. Praeger.

Schön, D. A. (1983). *The Reflective Practitioner: How Professionals Think in Action*. Basic Books.

Tetlock, P. E. (2005). *Expert Political Judgment: How Good Is It? How Can We Know?* Princeton University Press.

Tetlock, P. E., & Gardner, D. (2015). *Superforecasting: The Art and Science of Prediction*. Crown.