Assessment report

Candidate A · AI Automation Engineer (Project-Based)

> A junior-to-mid AI engineer with genuine but narrow depth in RAG systems and graph databases, honest about significant gaps in the automation tooling, CRM integration, and voice agent experience that are central to this role.

Interview recommendation

OUTCOME	OUTCOME	OUTCOME
Not Recommended to Interview ◀ THIS CANDIDATE	Recommended to Interview	Strongly Recommended to Interview

SCORE OVERVIEW

Technical Depth	3.0
Creative Reframing	2.0
Ethical Reasoning	2.0
Judgment Under Ambiguity	3.0
Intuition Under Scarcity	2.0
Overall	2.4 / 5
Cheating risk	Low

AT A GLANCE

- Narrow RAG depth
- Honest about gaps
- No automation tooling
- No voice-agent experience
- Weak role transfer
- Low cheating risk

Real but research-adjacent RAG/graph experience. Missing the core stack this role needs — n8n/Make/Zapier, CRM integration, end-to-end voice agents. Reasoning is pattern-level, not first-principles. Candid and low-risk, but a high-effort transfer.

Dimension scores — detail

Technical Depth — 3.0 / 5

“we were pairing semantic similarity retrieval with structured graph retrieval. Sorry, graph traversal, which gave us a lot more richer results than just pure vector approach alone”

Demonstrates genuine familiarity with RAG architecture and hybrid retrieval concepts. Can articulate the tradeoff between vector and graph approaches with some specificity. However, depth is uneven: the graph pruning solution was a borrowed time-threshold heuristic ('I have taken that hyperparameter of, say, 15 seconds from past projects... from the internet'), not derived from first principles. Hallucination mitigation answer required a long pause and landed at a fairly generic 'hybrid search' response before recovering with a more concrete constraint-validation example. No hands-on experience with n8n, Make, Zapier, CRM integrations, or voice agent platforms beyond standalone ElevenLabs TTS. Core automation tooling for this role is absent.

Creative Reframing — 2.0 / 5

“I think I'll be able to learn fairly quickly because that doesn't sound like a difficult tool to use”

When confronted with significant gaps (no CRM integration, no voice agent experience, no n8n/Make/Zapier), the candidate defaulted to confidence assertions rather than reframing the problem or proposing an alternative angle. No evidence of lateral thinking or novel problem decomposition. The clarifying-questions fallback in the RAG system was designed by senior engineers, not the candidate. Responses to novel probes were largely linear.

Ethical Reasoning 2.0 / 5

“those parts were mostly designed by the senior engineers on the team and I was more in charge of implementing and also performing audits”

No ethical reasoning was volunteered or demonstrated. The role's questions around compliance (outbound calling, opt-out) were not reached in the interview, and the candidate did not surface any ethical considerations unprompted. The one moment that could have touched on responsible AI design (constraint validation for a domain-specific bot) was framed purely as a technical quality problem, not an ethical one. Score capped at 2 due to absence of any substantive ethical signal.

Judgment Under Ambiguity 3.0 / 5

“The hardest query to handle is probably the ambiguous ones because not everyone communicates with the clearest intent. Sometimes the intent can sit between two domains and neither routing path is going to be clearly correct”

Candidate correctly identified ambiguity as the core challenge in intent routing and described a reasonable two-tier approach (keyword parsing for clear cases, LLM for edge cases, clarifying questions as final fallback). However, when pressed on how the clarifying question options were chosen, the candidate deferred to senior engineers. When asked how to derive a pruning threshold from first principles, the candidate admitted inability. Judgment is present at the pattern-recognition level but does not extend to first-principles reasoning under genuine uncertainty.

Intuition Under Scarcity 2.0 / 5

“Well, obviously I would want to know the exact input-output history of the users. So basically what kind of prompt the users have been giving and what kind of feedback, sorry, or output the bots are actually giving”

When asked for a gut-instinct first move on a broken customer service bot, the candidate gave a reasonable but generic answer (check the logs). Follow-up probing for the two or three most likely culprits produced a conditional, hedged response ('it really depends') rather than a prioritized intuition. The candidate did not demonstrate the ability to make a confident, reasoned call with limited information. The pause before the hallucination answer further suggests reliance on recall rather than intuitive pattern-matching.

Top excerpts

> we were pairing semantic similarity retrieval with structured graph retrieval... which gave us a lot more richer results than just pure vector approach alone. And second, I built the agentic workflow layer on top of that, specifically the intent recognition component

Most technically substantive claim in the interview. Establishes genuine familiarity with hybrid RAG architecture and agentic routing design.

> I have taken that hyperparameter of, say, 15 seconds from past projects, not my projects, but like past projects of a similar nature from the internet

Reveals the ceiling of independent technical reasoning. The candidate borrowed a heuristic rather than deriving it, and when pressed on first-principles derivation, could not articulate an approach.

> I have honestly never worked with such workflows. I think these are more web-related, while the ones I made are more, not web-related, but it simply is not within my tech stack

Direct confirmation of a critical gap for this role. n8n, Make, and Zapier are core to the job description and the candidate has zero hands-on experience.

> it's a combination of LLM generation and knowledge specific constraints that are written manually... the expert is going to write the high level guidelines. And then we take those guidelines and we allow another LLM to generate the constraints. But once the constraints are set, they are no longer variable

Most concrete and specific technical answer in the interview. Describes a real design pattern (human-in-the-loop constraint generation with LLM expansion) with clear architectural reasoning.

> No I haven't [built the full voice agent loop end to end]

Confirms absence of end-to-end voice agent experience, which is explicitly listed as a core requirement in the job description.

> I think I'll be able to learn them fairly quickly because that doesn't sound like a difficult tool to use

Candidate's response to a significant skill gap is confidence assertion without evidence. Contrasts with the more honest and specific answer given about LangChain/LangGraph learning curve.

Capability map

Transfer capability

Weak for this specific role. The candidate's demonstrated skills are concentrated in research-adjacent RAG and graph database work. The role requires production automation tooling (n8n, Make), CRM integration, voice agent platforms, and email pipeline construction — none of which the candidate has hands-on experience with. Transfer from academic/research AI engineering to commercial automation engineering is non-trivial and the candidate has not demonstrated the bridging skills or delivery track record that would make that transfer low-risk on a project basis.

Blind spots

- **Automation tooling (n8n, Make, Zapier):** zero hands-on experience; candidate's confidence that these are easy to learn is unsubstantiated.
- **CRM integrations (HubSpot, GoHighLevel, Zoho):** no experience with webhook-based pipelines or CRM connectors.
- **AI voice agents (Vapi, Retell, Bland AI):** no end-to-end voice agent experience; ElevenLabs used only for standalone TTS.
- **First-principles reasoning:** when asked to derive a pruning threshold from scratch, candidate could not articulate an approach.
- **Outbound calling compliance, email automation pipelines:** not discussed; no signals available.